

Clinsight Expert Directory — Assignment Feedback

Candidate: **Mahesh Kumar**

Overall outcome

Bugs correctly fixed: 7 / 8. Strong bug coverage and the most thorough written report of the three. However, the submission raises concerns about both process discipline and whether the work was done unaided (see note below), which we would want to resolve before progressing.

Bug-by-bug results

Assessed against the 8 planted bugs. Both planted red herrings (the “No rating yet” label and the // HACK seniority comment) were correctly left untouched.

Bug	Description	Outcome
B1	Case-sensitive search	Fixed
B2	Header shortlist count never updates	Fixed
B3	Form error doesn't clear after fix	Fixed
B4	Pagination repeats an item between pages	Fixed
B5	Sort-by-rating is ascending (worst first)	Fixed
B6	Specialty filter substring match	Missed
B7	Accepts bad email / negative years	Fixed
B8	Average rating inflated (off-by-one)	Fixed

What went well

- Caught and correctly fixed 7 of 8 bugs, including the subtle ones (off-by-one rating denominator, pagination overlap, ascending rating sort).
- By far the most professional FLOW_NOTES report: clear reproduction steps, expected vs. actual, and severity for every finding.
- Good judgement note: identified the back-navigation pagination issue and consciously documented it as out-of-scope rather than over-fixing.
- Both red herrings were left untouched.

Areas to improve

- Missed B6 (specialty filter substring match) — and explicitly recorded it as working correctly. Filtering “Cardiology” also returns “Pediatric Cardiology”.
- The brief asked for minimal, focused changes. Entire files were reformatted (single→double quotes, semicolons added everywhere), turning small fixes into a 300+ line diff that is hard to review.
- Added several things that were not asked for and, in one case, not actually a bug: shortlist persistence (the shortlist is intentionally session-only), duplicate-email checks, a name regex, and a max-experience cap.
- All work landed in a single bulk commit, so there is no visible testing/iteration trail.

Note on assignment conditions

We asked candidates not to use AI assistance for this task. Several characteristics of this submission are consistent with AI-generated work: the entire change landed as one large commit (no incremental testing trail), whole files were re-formatted rather than surgically edited, and the PR description and notes follow a highly templated, repetitive structure. We cannot prove this, but we think it is likely that AI was used, and we would want to discuss it directly before drawing conclusions about the underlying ability the assignment is meant to measure.